

## LOW-POWER DESIGN: OUTLINE

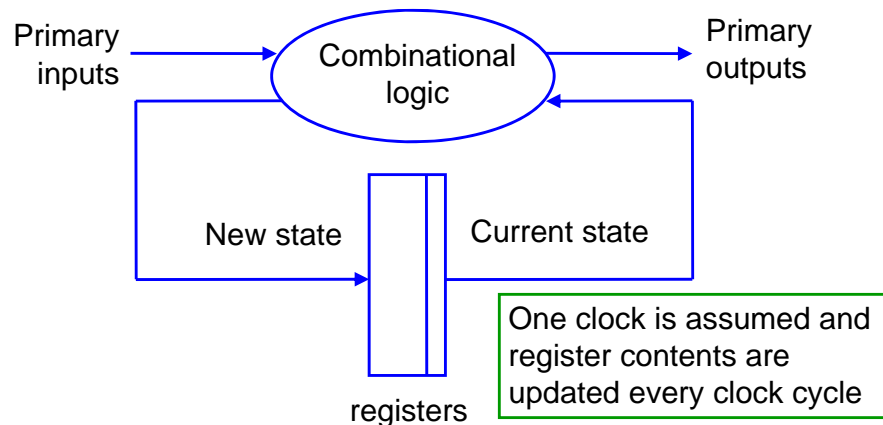
- Wrap-up register-transfer level design
- Implementation techniques: CMOS, standard cells, clock trees
- Sources of power dissipation
- Signal transition probabilities
- Power simulation and estimation techniques
- RT-level low-power design: with variable and constant voltage
- Architectural considerations
- Software-level control
- Asynchronous design
- Further reading

## RELATED COURSE

- 2<sup>nd</sup> quarter course:

*Energy-Efficient Embedded Systems (192130122)*

## REGISTER-TRANSFER (RT) VIEW OF HARDWARE




## REGISTER-TRANSFER MODEL IN PRACTICE

- An IC with some degree of complexity will have more than one clock.
- Register-transfer model is valid for each clock separately.
- One says that registers connected to the same clock belong to the same *clock domain*.
- One needs to take special care of signals crossing clock domains.


## CMOS TECHNOLOGY

- CMOS = *complementary metal-oxide-silicon*
- Building blocks: *n-type* and *p-type transistors*
- A transistor can be considered to act as a switch.

- A *p-type* transistor conducts current when its gate (G) is *low*.
- So, S=D when G=0; S is isolated from D when G=1.




- An *n-type* transistor conducts current when its gate is *high*.
- So, S=D when G=1; S is isolated from D when G=0.




## CMOS INVERTER

- The inverter:  $\text{out} = \text{not}(a)$ 
  - When  $a = 1$ , out is connected to  $V_{SS}$ , so  $\text{out} = 0$ .
  - When  $a = 0$ , out is connected to  $V_{DD}$ , so  $\text{out} = 1$ .
  - Ideally, no *static* current! ("static" means: stable input value)
  - Some *dynamic* current, see later.



## CMOS COMPLEX GATE

- A CMOS gate consists of two networks:
  - Network of p-type transistors for the connection to  $V_{DD}$
  - Network of n-type transistors for the connection to  $V_{SS}$
  - The two networks are complementary: a series connection in one is a parallel connection in the other.
  - Ideally, no static current!



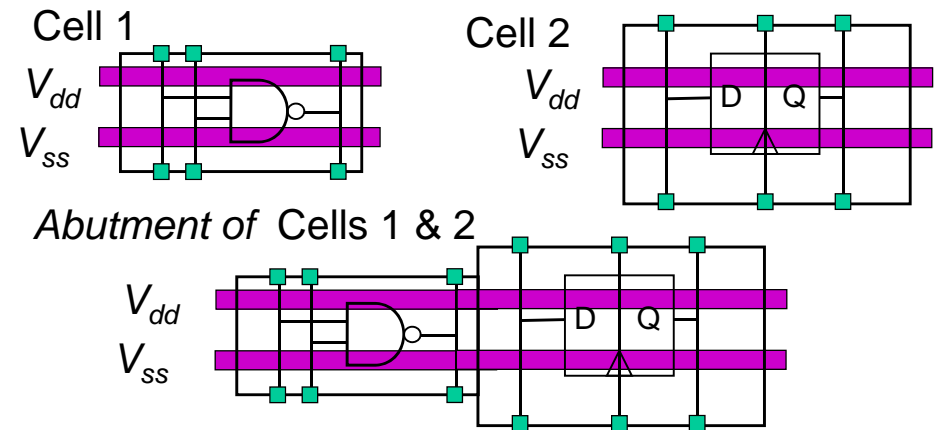
## STANDARD-CELL BASED DESIGN (1)

- The designer is given a *library* of elementary circuits called *cells*.
- The design of library cells (from transistors) is a specialized task and is done by a dedicated group or company.
- Library information consists of:
  - cell function and simulation models (delays!)
  - cell layouts.

## STANDARD-CELL BASED DESIGN (2)


- Designing amounts to delivering a *netlist* of cells for layout.
- The netlist is generally obtained by means of *logic synthesis*.
- The netlist is mapped on a layout by means of *placement* and *routing*.

## STANDARD-CELL LAYOUT



## TIME-AREA TRADE-OFF IN LOGIC SYNTHESIS


- In general, one can make a circuit faster by using more area.
- Example: a *carry-ripple adder* is small but relatively slow; a *carry-lookahead adder* is faster, but larger.
- In logic synthesis one constrains the clock period and the synthesis tool will choose the implementation that fits the constraint.



## CLOCK DISTRIBUTION

- Goals:
  - deliver clock to all memory elements with acceptable *skew* ("skew": difference between earliest and latest arrival of clock edge to a register);
  - deliver clock edges with acceptable sharpness.
- Clocking network design is one of the greatest challenges in the design of a large chip.


## H-TREE



Wire length  
from source to  
any endpoint  
is equal!


## CLOCK TREE

- In order to balance the delay from the clock source to the flip-flops, *clock trees* are used.
- In current-day practice clock trees are generated during layout as wiring delay is significant.



## SOURCES OF POWER DISSIPATION IN CMOS

- Charging and discharging nodes due to switching: *capacitive dissipation*. Energy loss:  
$$C \times V_{DD}^2$$
- A direct short-circuit path through both p-type and n-type transistors during switching: *short-circuit dissipation*
- Leakage



## POWER DISSIPATION EQUATION

- The power related to charging and discharging nodes is:

$$P = f \times C_{\text{eff}} \times V_{dd}^2$$


- Where:
  - $f$  : clock frequency
  - $C_{\text{eff}}$  : *effective* capacitance, average capacitance switched in one clock cycle
  - $V_{dd}$  : power-supply voltage

## WAYS TO REDUCE POWER

- Reduce clock frequency
- Reduce effective capacitance
- Reduce supply voltage:
  - Has a quadratic effect on power savings
  - Has a negative effect on switching delay
- Avoid unnecessary switching (glitching)

## GLITCHING: SWITCHING NOT CONTRIBUTING TO FUNCTION

- Example, using *unit-delay* simulation (1 time unit for each signal transition):



- Glitches *propagate* through the logic resulting in unnecessary power consumption

## SWITCHING ACTIVITY: TRANSITION PROBABILITY

- Probability of some signal  $s$  being logic '1':  $P_s$
- Then the signal transition probability for a signal that is *temporally uncorrelated*, becomes:

$$P_{tr} = 2P_s(1 - P_s)$$

- Examples:

$$P_s = \frac{1}{2} \rightarrow P_{tr} = \frac{1}{2} \quad P_s = \frac{1}{4} \rightarrow P_{tr} = \frac{3}{8}$$

## SWITCHING ACTIVITY FOR LOGIC GATES

- Probability of gate inputs being '1' can be translated to probability of outputs being one:

$$\text{Inverter: } P_{out} = 1 - P_{in}$$

$$\text{AND gate: } P_{out} = \prod_k P_{in_k}$$

$$\text{OR gate: } P_{out} = 1 - \prod_k (1 - P_{in_k})$$


- Assumption: inputs switch independently (often not true)

## TEMPORALLY CORRELATED SIGNALS

- One should be careful with the assumption that signals are temporally independent.
- A clock signal, for example, will always have 2 transitions per clock period, even when its probability of being '1' is 0.5.
- Consider the MSB of a sine wave sampled 20 times per period and encoded in 2's complement. Its probability of being '1' is 0.5. However, its transition probability is 0.1 (two transitions per period).

## SPATIALLY CORRELATED SIGNALS

- Signal having a common origin will be *spatially correlated*.
- In case of *reconvergent fanouts* independence is lost.



- MSBs in a bus, representing mostly small numbers, using 2's complement encoding, will also be spatially correlated.

## POWER-ESTIMATION TECHNIQUES (1)

- Circuit-level simulation
  - Use e.g. SPICE
  - All currents and voltages are known
  - Power can be computed accurately
  - Time consuming
- Gate-level simulation
  - Simulate standard-cell netlist
  - Keep track of switching activities
  - Combine with known power numbers for each cell
  - Glitches are taken into account

## POWER-ESTIMATION TECHNIQUES (2)

- Analysis-based estimation:
  - Propagation of probabilities
  - Correction for correlated signals
- Architecture-level estimation:
  - Based on characterization of building blocks, such as ALUs, memories, etc.
  - Combined with switching activity at the blocks' inputs.


## VOLTAGE REDUCTION

- Most effective
- Difficult to apply on individual designs:
  - Standard-cell library has been characterized for a specific voltage range.
- Even more difficult to use multiple voltages:
  - For the digital circuitry on the same IC
  - Within the same clock domain.
- Still, quite some research has been performed on designing with multiple voltages in the same clock domain
- Not supported by common logic synthesis tools

## AREA-POWER TRADE-OFF (1)


- Suppose a processor  $P_0$  with effective capacitance  $C_{\text{eff}}$  performs some task with given deadline operating at some frequency-voltage combination  $(f, V_{\text{dd}})$ .
- So, its power consumption is:

$$P = f \times C_{\text{eff}} \times V_{\text{dd}}^2$$



## AREA-POWER TRADE-OFF (2)

- Suppose now that two processors can be deployed to perform the task:
  - The effective capacitance will more than double (some overhead to distribute tasks).
  - The task's deadline can now be met by reducing the operating frequency to about half the original value.
  - No positive effect on power if voltage is kept constant.



## AREA-POWER TRADE-OFF (3)


- As the frequency was halved, the voltage can also be more or less halved without causing circuit malfunction (of course, the voltage cannot be reduced indefinitely).
- In the new situation:

$$\left. \begin{aligned} C'_{\text{eff}} &= 2.1 C_{\text{eff}} \\ f' &= 0.5 f \\ V'_{\text{dd}} &= 0.5 V_{\text{dd}} \end{aligned} \right\} \begin{aligned} P' &= 0.26 f \times C_{\text{eff}} \times V_{\text{dd}}^2 \\ &= 0.26 P \end{aligned}$$

Power reduction at the expense of area!

## COMPUTATIONAL MARGIN


- Critical path determines maximal clock period.
- Critical path may be reduced by increasing parallelism.
- Use the created *slack* or *computational margin* in the clock period, to reduce voltage.
- So, there may be power-area trade-off (if the effective capacitance has less than quadratic growth).



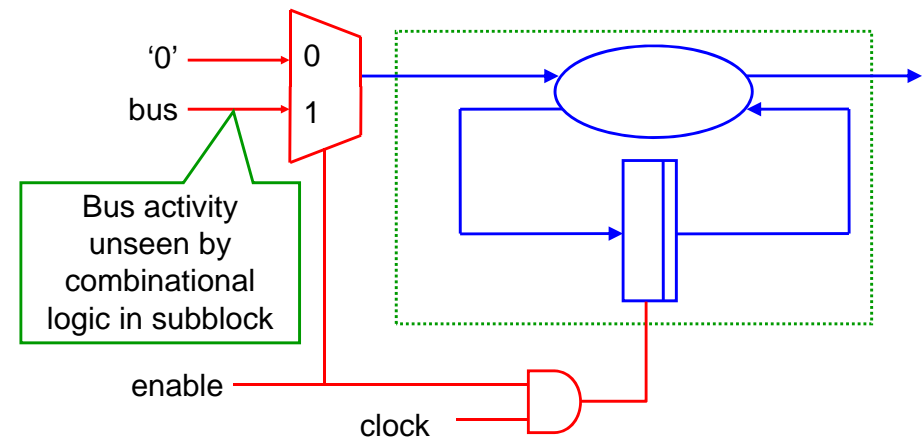
## CONSTANT-VOLTAGE RTL POWER REDUCTION

- Targeted to removing unnecessary switching including glitching.
- The easiest thing is to turn off entire block of a design when not in use:
  - Example: some blocks on a radio IC are only used for *transmitting* others only for *receiving*; the IC will probably do neither most of the time.
  - Implemented by *clock gating*

## BLOCK-LEVEL CLOCK GATING




## PREVENT COMBINATIONAL ACTIVITY






## PREVENT GLITCH PROPAGATION WITH REGISTERS (1)



## PREVENT GLITCH PROPAGATION WITH REGISTERS (2)


- One normally can't just add registers without modifying circuit functionality:
  - Pipelining* is allowed in acyclic data flow; it increases the *latency* (number of clock cycles for input to propagate to output).
  - Retiming* is always allowed (it does not change the number of registers in cyclic or acyclic path).




## INTRA-BLOCK CLOCK GATING (1)

- Registers do not always change value at each clock cycle; they sometimes keep their previous value (e.g. no "else" branch for new state computation in VHDL).
- In the latter case, the clock can be disabled.
- This saves power consumption in the clock tree and inside the register.
- This idea is encountered in many places, for example:
  - Theeuwes, F. and E. Seelen, *Power Reduction Through Clock Gating by Symbolic Manipulation*, VLSI '97, Gramado, Brazil, (1997).
- Also implemented in Synopsys tools.

## INTRA-BLOCK CLOCK GATING (2)



## INTRA-BLOCK CLOCK GATING (3)



## POWER GATING

- Clock gating preserves the circuit's state.
- For situations when batteries should last extremely long (several years), as for a node in a sensor network, clock gating is not enough. Leakage cannot be neglected.
- The solution is to turn off the power: *power gating*.
- Only a small subcircuit remains powered:
  - It periodically turns on and off the processor;
  - It stores the state information;
  - It may operate at a reduced clock frequency.

## ARCHITECTURAL CONSIDERATIONS

- Traditional processor has the *Von Neumann* architecture:
  - Program code and data in same memory.
  - This *bottleneck* not only affects performance but power as well (larger memories require more energy per access).
  - Time and energy is lost in memory transfers.
  - Multiplexing data and code streams destroys temporal correlation.
- Low-power architectures have *locality of reference*:
  - a high degree of parallelism
  - with local memories to avoid data transport across large distances.

## SOFTWARE-LEVEL LOW POWER DESIGN

- Compact code not only reduces the clock cycles needed for some computation, but also the switching activity and therefore the power consumption.
- Software (e.g. operating system that knows about the tasks to be performed) can control:
  - Voltage
  - Clock frequency
- Having dedicated co-processors for computation-intensive tasks will, in general, reduce power.
- Careful choice of sleep policy contributes to power reduction (wake-up time is relevant).

## ASYNCHRONOUS DESIGN

- The ultimate low-power implementation (given a technology).
- No clocks, no losses in clock trees, etc.
- Data stability is encoded in the data itself rather than guaranteed by clock edge.
- *Handshaking* is used to ensure that data transports are successful.
- There is only activity that is functionally meaningful.
- Computation speed adapts to requirements of the moment.
- Commercial designs exist, such as ARM processors.
- As well as design tool vendors: e.g. *Handshake Solutions* ☹️.

## FURTHER READING (1)

- Review papers:
  - Pedram, M., *Power Minimization in IC Design: Principles and Applications*, ACM Transactions on Design Automation of Electronic Systems, Vol.1(1), (January 1996).
  - Pedram, M. and A. Abdollahi, *Low-Power RT-Level Synthesis Techniques: A Tutorial*, IEE Proceedings on Computers and Digital Techniques, Vol.152(3), pp. 333-343, (May 2005).
  - Benini, L., G. De Micheli and E. Macii, *Designing Low-Power Circuits: Practical Recipes*, IEEE Circuits and Systems Magazine, Vol.1(1), pp. 6-25, (2001).

## FURTHER READING (2)

- Some books:
  - Chandrakasan, A.P. and R.W. Brodersen, *Low Power Digital CMOS Design*, Kluwer Academic Publishers, Boston, (1995).
  - Rabaey, J.M. and M. Pedram (Eds.), *Low Power Design Methodologies*, Kluwer Academic Publishers, Boston, (1995).
  - Benini, L. and G. De Micheli, *Dynamic Power Management, Design Techniques and CAD Tools*, Kluwer Academic Publishers, Boston, (1998).
  - Piguet, C. (Ed.), *Low-Power Electronics Design*, CRC Press, Boca Raton, (2005).

## FURTHER READING (3)

- Some more books:
  - Piguet, C. (Ed.), *Low-Power Processors and Systems on Chips*, Taylor and Francis, Boca Raton, (2006).
  - Otis, B. and J. Rabaey, *Ultra-Low Power Wireless Technologies for Sensor Networks*, Springer, (2007).
  - Mohanty, S.P., N. Ranganathan, E. Kougianos and P. Patra, *Low-Power High-Level Synthesis for Nanoscale CMOS Circuits*, Springer, (2008).



## FURTHER READING (4)

- A recent book by an expert in the field (on-line copy available through the UT Library):
  - Rabaey, J., *Low Power Design Essentials*, Springer, (2009).
  - <http://www.springerlink.com/content/978-0-387-71712-8>
- You are welcome to search in my database:
  - <http://www.bibix.nl/literature-db/search.html>